

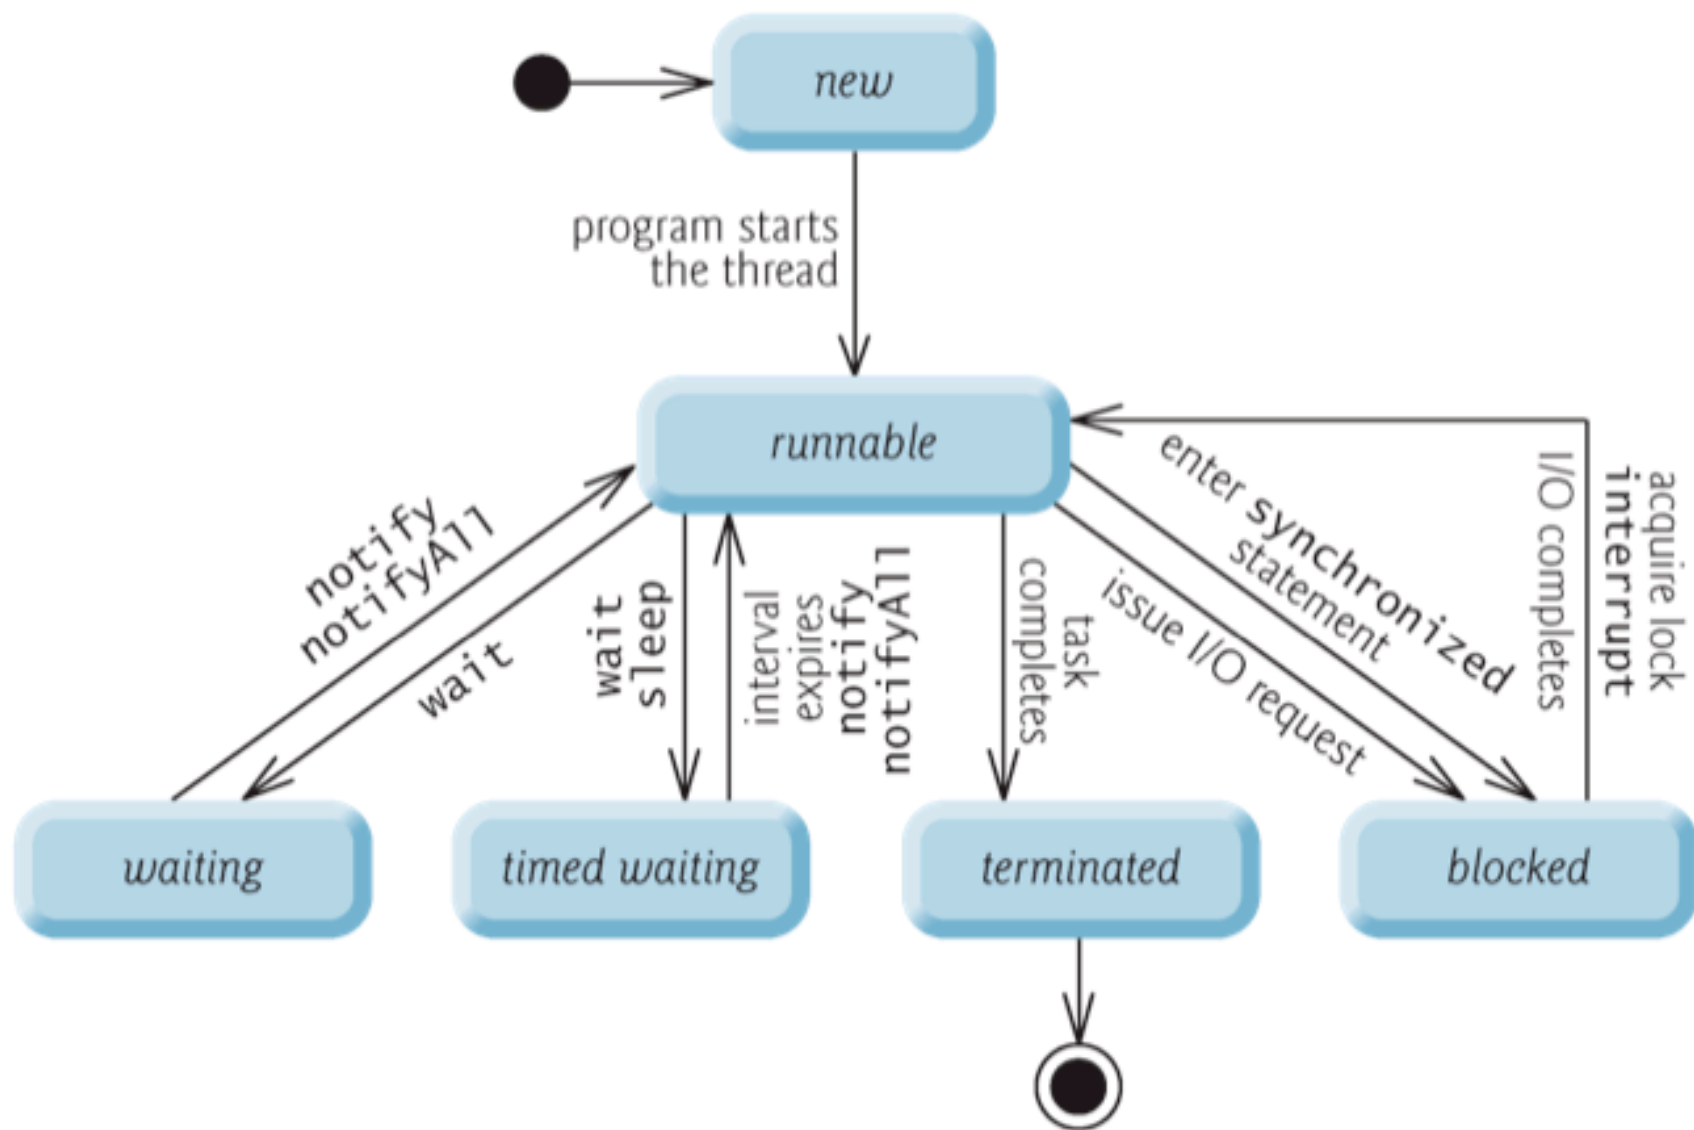
Multithreading

Introduction

- Java makes concurrency available to you through the language and APIs.
- You specify that an application contains separate **threads of execution**
 - each thread has its own **method-call stack** and program counter
 - can execute concurrently with other threads while sharing applicationwide resources such as memory with them.
- This capability is called **multithreading**.

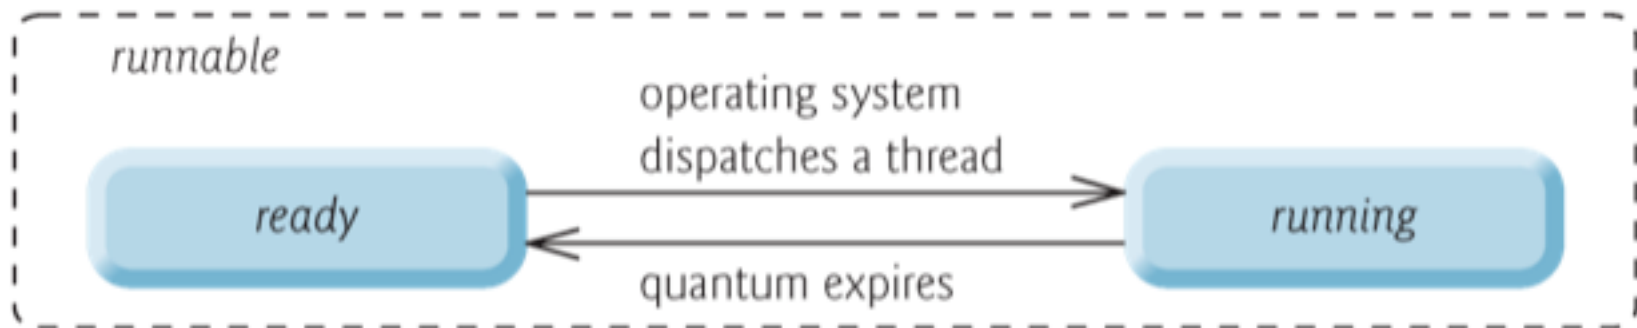
Thread States: Life Cycle of a Thread

- At any time, a thread is said to be in one of several **thread states**
- A new thread begins its life cycle in the **new state**.
- Remains there until started, which places it in the **runnable state**—**considered to be executing its task**.
- A *runnable* thread can transition to the *waiting* state while it waits for another thread to perform a task.
 - Transitions back to the runnable state only when another thread notifies it to continue executing.



Thread States: Life Cycle of a Thread (cont.)

- At the operating-system level, Java's *runnable* state typically encompasses two separate states



Thread States: Life Cycle of a Thread (cont.)

- Typically, each thread is given a a **quantum** or **timeslice** in which to perform its task.
- The process that an operating system uses to determine which thread to dispatch is called **thread scheduling**.

Thread States: Lifecycle of a Thread (cont.)

- **Thread priority**: helps determine the order in which threads are scheduled.
- Informally, higher-priority threads are more important to a program and should be allocated processor time before lower-priority threads.
- Thread priorities cannot guarantee the order in which threads execute.
- Do not explicitly create and use Thread objects to implement concurrency.
- Rather, use the **Executor** interface.

Thread States: Lifecycle of a Thread (cont.)

- Most operating systems support timeslicing, which enables threads of equal priority to share a processor.
- An operating system's **thread scheduler** determines which thread runs next.
- One simple thread-scheduler implementation keeps the highest-priority thread running at all times and, if there's more than one highest-priority thread, ensures that all such threads execute for a quantum each in **round-robin fashion**. This process continues until all threads run to completion.

Thread States: Lifecycle of a Thread (cont.)

- When a higher-priority thread enters the *ready* state, the operating system generally preempts the currently *running* thread (an operation known as **preemptive scheduling**).
- Higher-priority threads could postpone—possibly indefinitely—the execution of lower-priority threads.
- **Indefinite postponement** is sometimes referred to as **starvation**.
- Operating systems employ a technique called *aging* to prevent starvation.

Thread States: Lifecycle of a Thread (cont.)

- Java provides higher-level concurrency utilities to hide much of this complexity and make multithreaded programming less error prone.
- Thread priorities are used behind the scenes to interact with the operating system, but most programmers who use Java multithreading will not be concerned with setting and adjusting thread priorities.

Creating and Executing Threads in Java

- A **Runnable** object represents a “task” that can execute concurrently with other tasks.
- The **Runnable** interface declares the single method **run**, which contains the code that defines the task that a **Runnable** object should perform.
- When a thread executing a **Runnable** is created and started, the thread calls the **Runnable** object's **run** method, which executes in the new thread.

```
1. public class Worker implements Runnable
2. {
3.     public static void main (String[] args)
4.     {
5.         System.out.println("This is currently running on the main thread, " +
6.             "the id is: " + Thread.currentThread().getId());
7.         Worker worker = new Worker();
8.         Thread thread = new Thread(worker);
9.         thread.start();
10.    }

11.    @Override
12.    public void run()
13.    {
14.        System.out.println("This is currently running on a separate thread, " +
15.            "the id is: " + Thread.currentThread().getId());

16.    }
17. }
```

When this code is run, here's the output that I got:

- This is currently running on the main thread, the id is: 1
- This is currently running on a separate thread, the id is: 9

Creating and Executing Threads with Executor Framework

Creating and Executing Threads with Executor Framework (cont.)

- Recommended that you use the **Executor** interface to manage the execution of **Runnable** objects for you.
 - Typically creates and manages a group of threads called a **thread pool** to execute **Runnable**s.
- **Executors** can reuse existing threads and can improve performance by optimizing the number of threads.
- **Executor** method **execute** accepts a **Runnable** as an argument.
- An **Executor** assigns every **Runnable** to one of the available threads in the thread pool.
- If there are no available threads, the **Executor** creates a new thread or waits for a thread to become available.

```
1 // Fig. 26.3: PrintTask.java
2 // PrintTask class sleeps for a random time from 0 to 5 seconds
3 import java.util.Random;
4
5 public class PrintTask implements Runnable
6 {
7     private final int sleepTime; // random sleep time for thread
8     private final String taskName; // name of task
9     private final static Random generator = new Random();
10
11     // constructor
12     public PrintTask( String name )
13     {
14         taskName = name; // set task name
15
16         // pick random sleep time between 0 and 5 seconds
17         sleepTime = generator.nextInt( 5000 ); // milliseconds
18     } // end PrintTask constructor
19
```



```
20 // method run contains the code that a thread will execute
21 public void run()
22 {
23     try // put thread to sleep for sleepTime amount of time
24     {
25         System.out.printf( "%s going to sleep for %d milliseconds.\n",
26                             taskName, sleepTime );
27         Thread.sleep( sleepTime ); // put thread to sleep
28     } // end try
29     catch ( InterruptedException exception )
30     {
31         System.out.printf( "%s %s\n", taskName,
32                             "terminated prematurely due to interruption" );
33     } // end catch
34
35     // print task name
36     System.out.printf( "%s done sleeping\n", taskName );
37 } // end method run
38 } // end class PrintTask
```

```
1 // Fig. 26.4: TaskExecutor.java
2 // Using an ExecutorService to execute Runnable's.
3 import java.util.concurrent.Executors;
4 import java.util.concurrent.ExecutorService;
5
6 public class TaskExecutor
7 {
8     public static void main( String[] args )
9     {
10         // create and name each runnable
11         PrintTask task1 = new PrintTask( "task1" );
12         PrintTask task2 = new PrintTask( "task2" );
13         PrintTask task3 = new PrintTask( "task3" );
14
15         System.out.println( "Starting Executor" );
16
17         // create ExecutorService to manage threads
18         ExecutorService threadExecutor = Executors.newCachedThreadPool();
19
20         // start threads and place in runnable state
21         threadExecutor.execute( task1 ); // start task1
22         threadExecutor.execute( task2 ); // start task2
23         threadExecutor.execute( task3 ); // start task3
24     }
25 }
```

```
25         // shut down worker threads when their tasks complete
26         threadExecutor.shutdown();
27
28         System.out.println( "Tasks started, main ends.\n" );
29     } // end main
30 } // end class TaskExecutor
```

Starting Executor

Tasks started, main ends

```
task1 going to sleep for 4806 milliseconds
task2 going to sleep for 2513 milliseconds
task3 going to sleep for 1132 milliseconds
task3 done sleeping
task2 done sleeping
task1 done sleeping
```

Starting Executor

task1 going to sleep for 3161 milliseconds.

task3 going to sleep for 532 milliseconds.

task2 going to sleep for 3440 milliseconds.

Tasks started, main ends.

task3 done sleeping

task1 done sleeping

task2 done sleeping

Thread Synchronization

- When multiple threads share an object and it is modified by one or more of them, indeterminate results may occur, unless access to the shared object is managed properly.
- The problem can be solved by giving only one thread at a time *exclusive* access to code that manipulates the shared object.
 - During that time, other threads desiring to manipulate the object are kept waiting.
 - When the thread with exclusive access to the object finishes manipulating it, one of the threads that was waiting is allowed to proceed.

Thread Synchronization (cont.)

- This process, called **thread synchronization**, coordinates access to shared data by multiple concurrent threads.
 - Ensures that each thread accessing a shared object excludes all other threads from doing so simultaneously—this is called **mutual exclusion**.

Thread Synchronization (cont.)

- A common way to perform synchronization is to use Java's built-in **monitors**.
 - Every object has a monitor and a **monitor lock** (or **intrinsic lock**).
 - Can be held by a maximum of only one thread at any time.
 - A thread must acquire the lock before proceeding with the operation.
 - Other threads attempting to perform an operation that requires the same lock will be *blocked*.
- To specify that a thread must hold a monitor lock to execute a block of code, the code should be placed in a **synchronized statement**.
 - Said to be **guarded** by the monitor lock

Thread Synchronization (cont.)

- The **synchronized** statements are declared using the **synchronized** keyword:
 - **synchronized** (*object*)
 {
 statements
 } **// end synchronized statement**
- where *object* is the object whose monitor lock will be acquired
 - *object* is normally **this** if it's the object in which the **synchronized** statement appears.
- When a **synchronized** statement finishes executing, the object's monitor lock is released.
- Java also allows **synchronized methods**.

Unsynchronized Data Sharing

- A `SimpleArray` object in the example will be shared across multiple threads.
- Will enable those threads to place `int` values into `array`.
- Line 26 puts the thread that invokes `add` to sleep for a random interval from 0 to 499 milliseconds.
 - This is done to make the problems associated with unsynchronized access to shared data more obvious.

```
1 // Fig. 26.5: SimpleArray.java
2 // Class that manages an integer array to be shared by multiple threads.
3 import java.util.Arrays;
4 import java.util.Random;
5
6 public class SimpleArray // CAUTION: NOT THREAD SAFE!
7 {
8     private final int[] array; // the shared integer array
9     private int writeIndex = 0; // index of next element to be written
10    private final static Random generator = new Random();
11
12    // construct a SimpleArray of a given size
13    public SimpleArray( int size )
14    {
15        array = new int[ size ];
16    } // end constructor
17
```

```
18 // add a value to the shared array
19 public void add( int value )
20 {
21     int position = writeIndex; // store the write index
22
23     try
24     {
25         // put thread to sleep for 0-499 milliseconds
26         Thread.sleep( generator.nextInt( 500 ) );
27     } // end try
28     catch ( InterruptedException ex )
29     {
30         ex.printStackTrace();
31     } // end catch
32
```

```
33 // put value in the appropriate element
34 array[ position ] = value;
35 System.out.printf( "%s wrote %2d to element %d.\n",
36     Thread.currentThread().getName(), value, position );
37
38 ++writeIndex; // increment index of element to be written next
39 System.out.printf( "Next write index: %d\n", writeIndex );
40 } // end method add
41
42 // used for outputting the contents of the shared integer array
43 public String toString()
44 {
45     return "\nContents of SimpleArray:\n" + Arrays.toString( array );
46 } // end method toString
47 } // end class SimpleArray
```

Unsynchronized Data Sharing (cont.)

- Class `ArrayWriter` implements the interface `Runnable` to define a task for inserting values in a `SimpleArray` object.
- The task completes after three consecutive integers beginning with `startValue` are added to the `SimpleArray` object.

```
1 // Fig. 26.6: ArrayWriter.java
2 // Adds integers to an array shared with other Runnable's
3 import java.lang.Runnable;
4
5 public class ArrayWriter implements Runnable
6 {
7     private final SimpleArray sharedSimpleArray;
8     private final int startValue;
9
10    public ArrayWriter( int value, SimpleArray array )
11    {
12        startValue = value;
13        sharedSimpleArray = array;
14    } // end constructor
15
16    public void run()
17    {
18        for ( int i = startValue; i < startValue + 3; i++ )
19        {
20            sharedSimpleArray.add( i ); // add an element to the shared array
21        } // end for
22    } // end method run
23 } // end class ArrayWriter
```

```
1 // Fig 26.7: SharedArrayTest.java
2 // Executes two Runnable's to add elements to a shared SimpleArray.
3 import java.util.concurrent.Executors;
4 import java.util.concurrent.ExecutorService;
5 import java.util.concurrent.TimeUnit;
6
7 public class SharedArrayTest
8 {
9     public static void main( String[] arg )
10    {
11        // construct the shared object
12        SimpleArray sharedSimpleArray = new SimpleArray( 6 );
13
14        // create two tasks to write to the shared SimpleArray
15        ArrayWriter writer1 = new ArrayWriter( 1, sharedSimpleArray );
16        ArrayWriter writer2 = new ArrayWriter( 11, sharedSimpleArray );
17
18        // execute the tasks with an ExecutorService
19        ExecutorService executor = Executors.newCachedThreadPool();
20        executor.execute( writer1 );
21        executor.execute( writer2 );
22
23        executor.shutdown();
24    }
```

```
25     try
26     {
27         // wait 1 minute for both writers to finish executing
28         boolean tasksEnded = executor.awaitTermination(
29             1, TimeUnit.MINUTES );
30
31         if ( tasksEnded )
32             System.out.println( sharedSimpleArray ); // print contents
33         else
34             System.out.println(
35                 "Timed out while waiting for tasks to finish." );
36     } // end try
37     catch ( InterruptedException ex )
38     {
39         System.out.println(
40             "Interrupted while waiting for tasks to finish." );
41     } // end catch
42 } // end main
43 } // end class SharedArrayTest
```



```
pool-1-thread-1 wrote 1 to element 0.  
Next write index: 1  
pool-1-thread-1 wrote 2 to element 1.  
Next write index: 2  
pool-1-thread-1 wrote 3 to element 2.  
Next write index: 3  
pool-1-thread-2 wrote 11 to element 0.  
Next write index: 4  
pool-1-thread-2 wrote 12 to element 4.  
Next write index: 5  
pool-1-thread-2 wrote 13 to element 5.  
Next write index: 6
```

Contents of SimpleArray:
[11, 2, 3, 0, 12, 13]

First **pool-1-thread-1** wrote the value 1 to element 0. Later **pool-1-thread-2** wrote the value **11** to element 0, thus *overwriting* the previously stored value.

Synchronized Data Sharing—Making Operations Atomic

- The output errors can be attributed to the fact that the shared object, `SimpleArray`, is not **thread safe**.
- If one thread obtains the value of `writeIndex`, there is no guarantee that another thread cannot come along and increment `writeIndex` before the first thread has had a chance to place a value in the array.
- If this happens, the first thread will be writing to the array based on a **stale value** of `writeIndex`—a value that is no longer valid.

Synchronized Data Sharing—Making Operations Atomic (cont.)

- An **atomic operation** cannot be divided into smaller suboperations.
- Can simulate atomicity by ensuring that only one thread carries out the three operations at a time.
- Atomicity can be achieved using the **synchronized** keyword.

```
1 // Fig. 26.8: SimpleArray.java
2 // Class that manages an integer array to be shared by multiple
3 // threads with synchronization.
4 import java.util.Arrays;
5 import java.util.Random;
6
7 public class SimpleArray
8 {
9     private final int[] array; // the shared integer array
10    private int writeIndex = 0; // index of next element to be written
11    private final static Random generator = new Random();
12
13    // construct a SimpleArray of a given size
14    public SimpleArray( int size )
15    {
16        array = new int[ size ];
17    } // end constructor
18
```

```
19 // add a value to the shared array
20 public synchronized void add( int value )
21 {
22     int position = writeIndex; // store the write index
23
24     try
25     {
26         // put thread to sleep for 0-499 milliseconds
27         Thread.sleep( generator.nextInt( 500 ) );
28     } // end try
29     catch ( InterruptedException ex )
30     {
31         ex.printStackTrace();
32     } // end catch
33
```

```
34         // put value in the appropriate element
35         array[ position ] = value;
36         System.out.printf( "%s wrote %2d to element %d.\n",
37             Thread.currentThread().getName(), value, position );
38
39         ++writeIndex; // increment index of element to be written next
40         System.out.printf( "Next write index: %d\n", writeIndex );
41     } // end method add
42
43     // used for outputting the contents of the shared integer array
44     public String toString()
45     {
46         return "\nContents of SimpleArray:\n" + Arrays.toString( array );
47     } // end method toString
48 } // end class SimpleArray
```

```
pool-1-thread-1 wrote 1 to element 0.  
Next write index: 1  
pool-1-thread-2 wrote 11 to element 1.  
Next write index: 2  
pool-1-thread-2 wrote 12 to element 2.  
Next write index: 3  
pool-1-thread-2 wrote 13 to element 3.  
Next write index: 4  
pool-1-thread-1 wrote 2 to element 4.  
Next write index: 5  
pool-1-thread-1 wrote 3 to element 5.  
Next write index: 6
```

```
Contents of SimpleArray:  
1 11 12 13 2 3
```